Lifecare Hospital ER Simulation

A comprehensive **Discrete Event Simulation (DES)** model for analyzing Emergency Room performance and optimizing staffing decisions. This project implements a queue-based simulation system that models patient flow through triage and consultation stages.

Project Overview

This simulation helps hospital administrators understand:

- Patient flow through the ER system
- Bottleneck identification in triage vs. consultation
- Resource utilization of nurses and doctors
- Wait time optimization strategies
- Staffing recommendations based on data analysis

System Model

The simulation models a two-stage ER process:

```
Patient Arrival → Triage Queue → Triage Nurse → Consultation Queue → Doctor → Discharge
```

System Parameters

- Patient Arrivals: Exponential distribution (mean: 4 minutes)
- **Triage Service**: Uniform distribution (3-6 minutes)
- **Consultation Service**: Uniform distribution (5-15 minutes)
- Shift Duration: 8 hours (480 minutes)
- Resources: 2 triage nurses, 3 doctors (configurable)

Project Structure

```
project-2/
— er_simulation.py  # Core simulation engine
— scenario_analysis.py  # Comprehensive scenario testing
— demo.py  # Quick demo script
— test_simulation.py  # Verification and testing
— requirements.txt  # Python dependencies
— README.md  # This documentation
— PROJECT_SUMMARY.md  # Detailed project summary
```

Installation

1. Clone or download the project files

2. Install dependencies:

```
pip install -r requirements.txt
```

Usage

Quick Demo

Run a single simulation with baseline parameters:

```
python demo.py
```

Comprehensive Analysis

Run all scenarios and generate recommendations:

```
python scenario_analysis.py
```

Custom Simulation

```
from er_simulation import ERSimulation

# Create simulation with custom parameters
sim = ERSimulation(
    num_triage_nurses=3,
    num_doctors=4,
    shift_duration=600, # 10 hours
    arrival_mean=3.5 # 3.5 minutes between arrivals
)

# Run simulation
results = sim.run()
sim.print_results()
```

% Code Architecture

Core Classes

1. Event System

```
class EventType(Enum):
    PATIENT_ARRIVAL = "patient_arrival"
    TRIAGE_COMPLETE = "triage_complete"
```

```
CONSULTATION_COMPLETE = "consultation_complete"

@dataclass
class Event:
    time: float
    event_type: EventType
    patient_id: int

def __lt__(self, other):
    return self.time < other.time</pre>
```

Purpose: Manages simulation events using a priority queue for chronological processing.

Simple Explanation:

- **EventType**: This is like a list of different things that can happen in the hospital. We have three main events:
 - PATIENT_ARRIVAL: When a new patient shows up at the ER
 - TRIAGE_COMPLETE: When a nurse finishes checking a patient
 - CONSULTATION_COMPLETE: When a doctor finishes seeing a patient
- Event: This is like a note that says "something happened at a specific time". Each event has:
 - time: When this event happened (like 10:30 AM)
 - event_type: What kind of event it was (patient arrived, triage finished, etc.)
 - patient_id: Which patient this event is about (like patient #5)
- __lt__ method: This tells the computer how to sort events by time. It's like putting events in order from earliest to latest.

2. Patient Tracking

```
@dataclass
class Patient:
   id: int
    arrival time: float
    triage_start_time: Optional[float] = None
    triage end time: Optional[float] = None
    consultation_start_time: Optional[float] = None
    consultation_end_time: Optional[float] = None
    @property
    def triage_wait_time(self) -> float:
        if self.triage_start_time is None:
            return 0
        return self.triage_start_time - self.arrival_time
    @property
    def total wait time(self) -> float:
        return self.triage_wait_time + self.consultation_wait_time
```

Purpose: Tracks individual patient journey and calculates wait/service times.

Simple Explanation:

• Patient: This is like a patient's medical record that follows them through the ER. It keeps track of:

```
o id: Patient's number (like patient #1, #2, etc.)
```

- o arrival_time: When they first walked into the ER
- triage_start_time: When a nurse started checking them
- triage_end_time: When the nurse finished checking them
- consultation_start_time: When a doctor started seeing them
- o consultation_end_time: When the doctor finished seeing them
- triage_wait_time: This calculates how long the patient had to wait before a nurse could see them. It's like: "When did the nurse start seeing them?" minus "When did they arrive?"
- total_wait_time: This adds up all the waiting time both waiting for the nurse and waiting for the
 doctor.

Think of it like a stopwatch that starts when the patient arrives and keeps track of every step of their journey through the ER.

3. Main Simulation Engine

```
class ERSimulation:
   def __init__(self,
                 num_triage_nurses: int = 2,
                 num doctors: int = 3,
                 shift_duration: float = 480.0,
                 arrival_mean: float = 4.0,
                 triage min: float = 3.0,
                 triage_max: float = 6.0,
                 consultation_min: float = 5.0,
                 consultation max: float = 15.0):
        # System parameters
        self.num_triage_nurses = num_triage_nurses
        self.num doctors = num doctors
        self.shift_duration = shift_duration
        # Simulation state
        self.current_time = 0.0
        self.events = [] # Priority queue
        self.patients = {} # Dictionary of all patients
        # Resource states
        self.available_triage_nurses = num_triage_nurses
        self.available doctors = num doctors
        # Queues
```

```
self.triage_queue = []
self.consultation_queue = []
# Statistics tracking
self.stats = {
    'total_patients': 0,
    'completed_patients': 0,
    'triage_wait_times': [],
    'consultation_wait_times': [],
    'total_wait_times': [],
    'triage_service_times': [],
    'consultation_service_times': [],
    'triage_queue_lengths': [],
    'consultation_queue_lengths': [],
    'triage_utilization': [],
    'doctor_utilization': [],
    'time_points': []
}
```

Simple Explanation: This is like setting up a virtual hospital ER. Think of it as creating a digital version of the emergency room with:

System Setup (like hospital planning):

- num_triage_nurses: How many nurses are working (default: 2)
- num_doctors: How many doctors are working (default: 3)
- shift_duration: How long the shift lasts in minutes (default: 480 = 8 hours)
- arrival_mean: How often patients arrive (default: every 4 minutes on average)
- triage_min/max: How long triage takes (3-6 minutes)
- consultation_min/max: How long doctor consultation takes (5-15 minutes)

Hospital State (like the current situation):

- current time: What time it is in the simulation (starts at 0)
- events: A list of things that will happen (like a schedule)
- patients: A list of all patients in the system

Available Staff (like who's free to work):

- available triage nurses: How many nurses are not busy right now
- available_doctors: How many doctors are not busy right now

Waiting Lines (like the waiting room):

- triage_queue: Patients waiting to see a nurse
- consultation queue: Patients waiting to see a doctor

Statistics (like keeping track of everything):

- total_patients: How many patients have come in total
- completed_patients: How many patients have finished their treatment
- triage_wait_times: How long each patient waited for a nurse

- consultation wait times: How long each patient waited for a doctor
- And many more tracking items...

Key Methods

Event Scheduling

```
def _schedule_next_arrival(self):
    """Schedule the next patient arrival using exponential distribution"""
    if self.current_time < self.shift_duration:
        inter_arrival_time = random.expovariate(1.0 / self.arrival_mean)
        arrival_time = self.current_time + inter_arrival_time

    if arrival_time <= self.shift_duration:
        event = Event(arrival_time, EventType.PATIENT_ARRIVAL,
    self.patient_id_counter)
        heapq.heappush(self.events, event)</pre>
```

Simple Explanation: This function is like a receptionist who decides when the next patient will arrive. Here's what it does:

- 1. Check if shift is still going: if self.current_time < self.shift_duration</pre>
 - Only schedule new patients if the hospital is still open
- 2. Calculate arrival time: inter_arrival_time = random.expovariate(1.0 / self.arrival_mean)
 - This uses a special math formula (exponential distribution) to randomly decide when the next patient will arrive
 - If arrival_mean = 4, patients arrive on average every 4 minutes, but sometimes it might be 2 minutes, sometimes 6 minutes

```
3. Create the arrival event: event = Event(arrival_time, EventType.PATIENT_ARRIVAL,
    self.patient_id_counter)
```

- This creates a note saying "Patient #X will arrive at time Y"
- 4. **Add to schedule**: heapq.heappush(self.events, event)
 - This puts the event in a special list that keeps everything in time order

Real-world example: If it's 10:00 AM and patients arrive every 4 minutes on average, this might schedule the next patient to arrive at 10:03 AM, then the next at 10:07 AM, etc.

Service Time Generation

```
def _generate_triage_time(self) -> float:
    """Generate triage service time using uniform distribution"""
    return random.uniform(self.triage_min, self.triage_max)
```

```
def _generate_consultation_time(self) -> float:
    """Generate consultation service time using uniform distribution"""
    return random.uniform(self.consultation_min, self.consultation_max)
```

Simple Explanation: These functions are like timers that decide how long each medical procedure will take. In real life, some patients take longer to treat than others.

```
_generate_triage_time():
```

- This decides how long a nurse will spend checking a patient
- Uses random.uniform(self.triage_min, self.triage_max) which means:
 - If triage_min = 3 and triage_max = 6
 - The nurse might spend 3 minutes, 4 minutes, 5 minutes, or 6 minutes
 - Each time is equally likely (uniform distribution)

```
_generate_consultation_time():
```

- This decides how long a doctor will spend with a patient
- Uses random.uniform(self.consultation_min, self.consultation_max) which means:
 - If consultation_min = 5 and consultation_max = 15
 - The doctor might spend 5 minutes, 8 minutes, 12 minutes, or 15 minutes
 - Each time is equally likely

Real-world example:

- A simple case might take 3 minutes for triage and 5 minutes for consultation
- A complex case might take 6 minutes for triage and 15 minutes for consultation
- The simulation randomly picks times within these ranges to make it realistic

Patient Flow Management

```
def _handle_patient_arrival(self, patient_id: int):
    """Handle a new patient arrival"""
    patient = Patient(id=patient_id, arrival_time=self.current_time)
    self.patients[patient_id] = patient
    self.stats['total_patients'] += 1

# Try to start triage immediately
    if self.available_triage_nurses > 0:
        self._start_triage(patient)
    else:
        # Add to triage queue
        self.triage_queue.append(patient)

# Schedule next arrival
    self._schedule_next_arrival()
```

Simple Explanation: This function is like a receptionist who greets each new patient and decides what happens next. Here's the step-by-step process:

- 1. Create patient record: patient = Patient(id=patient_id, arrival_time=self.current_time)
 - This creates a new patient file with their ID and arrival time
 - Like filling out a new patient form
- 2. Add to patient list: self.patients[patient_id] = patient
 - This puts the patient in the hospital's master list
 - · Like adding them to the hospital database
- 3. Update statistics: self.stats['total_patients'] += 1
 - This increases the count of total patients who have arrived
 - Like updating the daily patient count
- 4. Check if nurse is available: if self.available_triage_nurses > 0:
 - This checks if there's a free nurse to see the patient immediately
 - Like looking around to see if any nurses are free
- 5. If nurse is free: self._start_triage(patient)
 - Start the triage process right away
 - Like immediately calling a nurse to see the patient
- 6. If no nurse is free: self.triage_queue.append(patient)
 - Put the patient in the waiting line
 - Like telling the patient "Please take a seat, a nurse will be with you shortly"
- 7. Schedule next patient: self._schedule_next_arrival()
 - Plan when the next patient will arrive
 - Like the receptionist looking at their schedule

Real-world example:

- Patient #5 arrives at 10:30 AM
- Receptionist checks: "Is there a free nurse?"
- If yes: "Nurse Sarah, please see patient #5"
- If no: "Patient #5, please wait in the triage waiting area"

Resource Allocation

```
def _start_triage(self, patient: Patient):
    """Start triage for a patient"""
    self.available_triage_nurses -= 1
    patient.triage_start_time = self.current_time
```

```
# Schedule triage completion
triage_time = self._generate_triage_time()
triage_end_time = self.current_time + triage_time

event = Event(triage_end_time, EventType.TRIAGE_COMPLETE, patient.id)
heapq.heappush(self.events, event)
```

Simple Explanation: This function is like a nurse manager who assigns a nurse to a patient and schedules when they'll finish. Here's what happens:

- 1. Assign a nurse: self.available_triage_nurses -= 1
 - This marks one nurse as "busy" (reduces the count of available nurses)
 - Like saying "Nurse Sarah is now busy with a patient"
- 2. Record start time: patient.triage_start_time = self.current_time
 - This writes down exactly when the nurse started seeing the patient
 - Like clocking in when the nurse starts the triage
- 3. Calculate how long it will take: triage_time = self._generate_triage_time()
 - This randomly decides how long this specific triage will take (3-6 minutes)
 - Like estimating "This patient will take about 4 minutes to check"
- 4. Calculate finish time: triage_end_time = self.current_time + triage_time
 - This calculates when the nurse will finish
 - Like saying "If it's 10:30 AM now and it takes 4 minutes, they'll finish at 10:34 AM"
- 5. Schedule the completion: event = Event(triage_end_time, EventType.TRIAGE_COMPLETE,
 patient.id)
 - This creates a reminder that says "At 10:34 AM, nurse will finish with patient #5"
 - Like putting an appointment in the calendar
- 6. Add to schedule: heapq.heappush(self.events, event)
 - This puts the completion event in the hospital's master schedule
 - Like adding it to the daily planner

Real-world example:

- It's 10:30 AM
- Nurse Sarah starts seeing patient #5
- The system randomly decides this triage will take 4 minutes
- It schedules a reminder for 10:34 AM: "Nurse Sarah will finish with patient #5"

Main Simulation Loop

```
def run(self):
   """Run the simulation"""
   print(f"Starting ER simulation...")
   print(f"Parameters: {self.num_triage_nurses} nurses, {self.num_doctors}
doctors")
   print(f"Shift duration: {self.shift_duration} minutes")
   while self.events and self.current_time < self.shift_duration:
       # Get next event
       event = heapq.heappop(self.events)
       self.current time = event.time
       # Update statistics
        self._update_statistics()
       # Handle event
       if event.event_type == EventType.PATIENT_ARRIVAL:
            self._handle_patient_arrival(event.patient_id)
        elif event.event_type == EventType.TRIAGE_COMPLETE:
            self._handle_triage_complete(event.patient_id)
        elif event.event_type == EventType.CONSULTATION_COMPLETE:
            self._handle_consultation_complete(event.patient_id)
   print(f"Simulation completed at time {self.current_time:.2f} minutes")
   return self.get_results()
```

Simple Explanation: This is the main "brain" of the simulation - like a hospital manager who runs the entire ER shift. It's like a clock that ticks forward and handles everything that happens. Here's how it works:

1. Start the shift:

- Print the setup information (how many nurses, doctors, shift length)
- Like a manager starting their shift and reviewing the day's plan
- 2. Main loop: while self.events and self.current_time < self.shift_duration:</pre>
 - Keep going as long as there are events to handle AND the shift isn't over
 - Like saying "Keep working until there's nothing left to do or the shift ends"
- 3. Get the next event: event = heapq.heappop(self.events)
 - Take the next scheduled event from the list (the earliest one)
 - Like looking at your calendar and seeing "What's the next thing I need to do?"
- 4. Update the clock: self.current_time = event.time
 - Move the simulation time forward to when this event happens
 - Like saying "It's now 10:34 AM" (when the event is scheduled)
- 5. Update statistics: self. update statistics()
 - Record what's happening right now (queue lengths, staff busy/free)

- Like taking a snapshot of the current situation
- 6. Handle the event: The big decision tree
 - If it's a patient arrival: self._handle_patient_arrival(event.patient_id)
 - Like a receptionist greeting a new patient
 - If it's triage completion: self._handle_triage_complete(event.patient_id)
 - Like a nurse finishing with a patient and sending them to the doctor
 - o If it's consultation completion:

```
self._handle_consultation_complete(event.patient_id)
```

- Like a doctor finishing with a patient and discharging them
- 7. End the shift: When the loop ends, print completion message and return results

Real-world example:

- It's 10:30 AM: "Next event: Patient #5 arrives"
- Handle patient arrival
- It's 10:34 AM: "Next event: Nurse finishes with patient #5"
- Handle triage completion
- It's 10:40 AM: "Next event: Patient #6 arrives"
- And so on until the shift ends...

■ Performance Metrics

Wait Times

- Triage Wait: Time from arrival to triage start
- Consultation Wait: Time from triage completion to consultation start
- **Total Wait**: Sum of triage and consultation wait times

Simple Explanation: Think of these like waiting in line at different places:

Triage Wait: How long you wait before a nurse sees you

- Like waiting in line at the reception desk
- "I arrived at 10:00 AM, but the nurse didn't start seeing me until 10:05 AM"
- Wait time = 5 minutes

Consultation Wait: How long you wait before a doctor sees you

- · Like waiting in the doctor's waiting room after the nurse is done
- "The nurse finished at 10:08 AM, but the doctor didn't start seeing me until 10:15 AM"
- Wait time = 7 minutes

Total Wait: All the waiting time combined

• "I waited 5 minutes for the nurse + 7 minutes for the doctor = 12 minutes total"

Service Times

• Triage Service: Time from triage start to completion

- **Consultation Service**: Time from consultation start to completion
- **Total Service**: Sum of triage and consultation service times

Simple Explanation: These are the actual treatment times - how long the medical staff spend with each patient:

Triage Service: How long the nurse spends checking the patient

- Like the actual time the nurse is examining you
- "The nurse started checking me at 10:05 AM and finished at 10:08 AM"
- Service time = 3 minutes

Consultation Service: How long the doctor spends with the patient

- Like the actual time the doctor is treating you
- "The doctor started seeing me at 10:15 AM and finished at 10:25 AM"
- Service time = 10 minutes

Total Service: All the treatment time combined

"Nurse spent 3 minutes + doctor spent 10 minutes = 13 minutes total treatment time"

Queue Analysis

- Triage Queue Length: Number of patients waiting for triage
- Consultation Queue Length: Number of patients waiting for consultation
- Maximum Queue Lengths: Peak queue sizes during simulation

Simple Explanation: These tell you how crowded the waiting areas are:

Triage Queue Length: How many people are waiting to see a nurse

- Like counting people in the reception area
- "There are 3 people waiting to see a nurse right now"

Consultation Queue Length: How many people are waiting to see a doctor

- Like counting people in the doctor's waiting room
- "There are 5 people waiting to see a doctor right now"

Maximum Queue Lengths: The busiest the waiting areas ever got

- Like remembering "The reception area was most crowded with 8 people waiting"
- This helps understand peak demand periods

Resource Utilization

- Triage Nurse Utilization: Percentage of time nurses are busy
- **Doctor Utilization**: Percentage of time doctors are busy

Simple Explanation: These tell you how busy the medical staff are:

Triage Nurse Utilization: How much of their time nurses spend working vs. waiting

- Like saying "Nurses are busy 80% of the time, free 20% of the time"
- 80% utilization means nurses are working most of the time
- 20% utilization means nurses are often free and waiting for patients

Doctor Utilization: How much of their time doctors spend working vs. waiting

- Like saying "Doctors are busy 90% of the time, free 10% of the time"
- High utilization (90%) means doctors are very busy
- Low utilization (30%) means doctors often have free time

What this means:

- High utilization (80-90%): Staff are working hard, but patients might wait longer
- Low utilization (20-30%): Staff have free time, but might be overstaffed
- Balanced utilization (50-70%): Good balance between efficiency and patient care

Throughput

- Patients per Hour: Average patients completed per hour
- Total Patients Served: Number of patients who completed the entire process

Simple Explanation: These tell you how many patients the ER can handle:

Patients per Hour: How many patients complete their entire treatment per hour

- Like saying "The ER can treat 15 patients per hour on average"
- This measures the ER's capacity and efficiency
- Higher numbers mean the ER is more efficient

Total Patients Served: How many patients completed their full treatment during the shift

- Like saying "During the 8-hour shift, we treated 120 patients"
- This shows the total workload handled

What this means:

- **High throughput** (15+ patients/hour): ER is efficient and can handle many patients
- Low throughput (5-8 patients/hour): ER might be understaffed or inefficient
- Balanced throughput (10-15 patients/hour): Good balance of quality care and efficiency

Scenario Analysis

Staffing Configurations

The system tests various combinations:

```
staffing_configs = [
    (1, 2, "1 nurse, 2 doctors"),
    (1, 3, "1 nurse, 3 doctors"),
    (2, 2, "2 nurses, 2 doctors"),
    (2, 3, "2 nurses, 3 doctors"),
    (2, 4, "2 nurses, 4 doctors"),
```

```
(3, 3, "3 nurses, 3 doctors"),
(3, 4, "3 nurses, 4 doctors"),
(3, 5, "3 nurses, 5 doctors")
]
```

Shift Duration Testing

```
durations = [240, 360, 480, 600, 720] # 4, 6, 8, 10, 12 hours
```

Arrival Rate Scenarios

```
arrival_means = [2, 3, 4, 5, 6] # minutes between arrivals
```

Service Time Variations

```
service_configs = [
    ((2, 4), (4, 12), "Faster service times"),
    ((3, 6), (5, 15), "Baseline service times"),
    ((4, 8), (6, 18), "Slower service times"),
    ((5, 10), (8, 20), "Much slower service times")
]
```

Output Files

The comprehensive analysis generates:

Data Files

• scenario_comparison.csv: Detailed results table with all metrics

Visualization Charts

- wait_times_comparison.png: Wait time comparison across scenarios
- utilization_comparison.png: Resource utilization comparison
- queue lengths.png: Queue length comparison

Example Results Table

Scenario	Total Patients	Completed Patients	Patients/Hour	Avg Total Wait	Triage Utilization	Doctor Utilization
	Patients	ratients		(min)	Othization	Othization

Scenario	Total Patients	Completed Patients	Patients/Hour	Avg Total Wait (min)	Triage Utilization	Doctor Utilization
Baseline (2 nurses, 3 doctors)	121	119	14.88	9.19	67.1%	87.2%
1 nurse, 2 doctors	98	85	10.62	10.56	86.6%	87.2%
2 nurses, 4 doctors	120	111	13.88	1.36	64.0%	65.7%

***** Key Findings

Baseline Performance (2 nurses, 3 doctors, 8-hour shift):

• Average Total Wait Time: 9.19 minutes

Patients per Hour: 14.88 patients
Triage Nurse Utilization: 67.1%

• **Doctor Utilization**: 87.2%

• Maximum Queue Lengths: 5 (triage), 10 (consultation)

Optimal Configurations:

• Lowest Wait Time: 2 nurses, 4 doctors (1.36 min)

• Highest Throughput: 2 nurses, 3 doctors (14.88 patients/hour)

Best Resource Balance: 3 nurses, 3 doctors (42.4% nurse, 76.9% doctor utilization)

Bottleneck Analysis:

• **Primary Bottleneck**: Doctors (consultation wait time > triage wait time)

Recommendation: Add more doctors to reduce consultation queues

% Technical Implementation Details

Discrete Event Simulation

- Event-driven architecture using priority queues
- Realistic probability distributions for arrivals and service times
- Accurate tracking of patient flow through the system

Data Structures

- Patient: Tracks individual patient journey and timing
- Event: Represents simulation events with chronological ordering
- **ERSimulation**: Main simulation engine with comprehensive statistics

Statistics Collection

- **Real-time tracking** of queue lengths and resource utilization
- Comprehensive performance metrics calculation
- Scenario comparison and analysis capabilities

Requirements Fulfilled

✓ Queue-Based ER Process Simulation (30 marks)

- Complete two-stage ER workflow simulation
- Realistic patient arrival and service time distributions
- Proper queue management and resource allocation

✓ Performance Metrics Tracking

- Average patient wait times at each stage
- Resource utilization for nurses and doctors
- Queue lengths and throughput analysis
- Patients served per shift tracking

Scenario Testing Capabilities

- Different staffing configurations (1-3 nurses, 2-5 doctors)
- Various shift durations (4-12 hours)
- Different arrival rates and service time distributions
- Comprehensive comparison and analysis

☑ Recommendations Generated

- · Bottleneck identification and analysis
- · Optimal staffing recommendations
- Data-driven improvement suggestions
- Resource allocation strategies

Educational Value

This project demonstrates:

- Discrete Event Simulation principles and implementation
- Queue theory applications in healthcare
- Resource allocation optimization techniques
- Performance analysis and metrics calculation
- Data-driven decision making in operational management
- Python programming with object-oriented design
- Statistical analysis and visualization techniques

Future Enhancements

The simulation can be extended with:

Patient priority levels (urgent vs. non-urgent)

- Different triage categories and routing logic
- Real-time monitoring capabilities
- Cost analysis features
- Multiple shift scheduling and dynamic staffing
- Advanced queue management strategies

Example Usage

Running a Quick Test

python demo.py

Expected Output:

```
LIFECARE HOSPITAL ER SIMULATION DEMO
_____
Starting ER simulation...
Parameters: 2 nurses, 3 doctors
Shift duration: 480.0 minutes
Simulation completed at time 480.84 minutes
______
ER SIMULATION RESULTS
_____
Total Patients Arrived: 121
Patients Completed: 119
Patients per Hour: 14.88
WAIT TIMES (minutes):
 Average Triage Wait: 1.41
 Average Consultation Wait: 7.80
 Average Total Wait: 9.19
SERVICE TIMES (minutes):
 Average Triage Service: 4.59
 Average Consultation Service: 10.12
 Average Total Service: 14.72
QUEUE LENGTHS:
 Maximum Triage Queue: 5
 Maximum Consultation Queue: 10
RESOURCE UTILIZATION:
 Average Triage Nurse Utilization: 67.08%
 Average Doctor Utilization: 87.22%
_____
```

How to Read These Results:

Patient Flow:

- 121 patients arrived during the 8-hour shift
- 119 patients completed their full treatment (2 left without finishing)
- 14.88 patients per hour the ER treated about 15 patients every hour

Wait Times (Patient Experience):

- 1.41 minutes average wait for a nurse patients usually get seen quickly by nurses
- 7.80 minutes average wait for a doctor patients wait longer for doctors
- 9.19 minutes total wait time from arrival to starting treatment

Service Times (Treatment Duration):

- 4.59 minutes average nurse triage time nurses spend about 4-5 minutes checking patients
- 10.12 minutes average doctor consultation time doctors spend about 10 minutes with patients
- 14.72 minutes total treatment time from start to finish of medical care

Queue Analysis (Crowding):

- Maximum 5 patients waiting for nurses at any time
- Maximum 10 patients waiting for doctors at any time
- The doctor waiting area gets more crowded than the nurse area

Staff Utilization (Efficiency):

- 67.08% nurse utilization nurses are busy about 2/3 of the time
- 87.22% doctor utilization doctors are busy almost 90% of the time
- Doctors are working harder than nurses in this setup

What This Means:

- The ER is running efficiently (treating 15 patients/hour)
- Patients don't wait too long (under 10 minutes total wait)
- Doctors are very busy (87% utilization) might need more doctors
- Nurses have some free time (67% utilization) might have enough nurses

Running Comprehensive Analysis

python scenario_analysis.py

This will generate:

- Detailed comparison tables
- Visualization charts
- CSV export of results
- Data-driven recommendations



This project is created for educational purposes as part of CS401 coursework.

S Contributing

This simulation provides a solid foundation for healthcare operations research and can be extended with additional features like:

- Patient prioritization systems
- Multi-shift scheduling
- Cost-benefit analysis
- Real-time monitoring dashboards
- Integration with hospital information systems

Note: This simulation provides valuable insights for hospital administrators to optimize ER operations and improve patient satisfaction through reduced wait times and better resource utilization.